Notes for New Constructions of DMPF

tbd

ABSTRACT

tbd.

CCS CONCEPTS

• **Theory of computation → Cryptographic primitives**.

KEYWORDS

tbd

1 INTRODUCTION

tbd

2 PRELIMINARY

2.1 Basic Notations

Point and multi-point functions. Given a domain size N and Abelian group $\mathbb{G}$, a *point function* $f_{\alpha,\beta} : [N] \to \mathbb{G}$ for $\alpha \in [N]$ and $\beta \in \mathbb{G}$ evaluates to β on input α and to $0 \in \mathbb{G}$ on all other inputs. We denote by $\hat{f}_{\alpha,\beta} = (N, \hat{\mathbb{G}}, \alpha, \beta)$ the representation of such a point function. A *t-point function* $f_{A,B} : [N] \to \mathbb{G}$ for $A = (\alpha_1, \cdots \alpha_t) \in [N]^t$ and $B = (\beta_1, \cdots, \beta_t) \in \mathbb{G}^t$ evaluates to β_i on input α_i for $1 \le i \le t$ and to 0 on all other inputs. Denote $\hat{f}_{A,B}(N, \hat{\mathbb{G}}, t, A, B)$ the representation of such a t-point function. Call the collection of all t-point functions for all t multi-point functions.

Enote: MPF. Also representation of groups.

2.2 Distributed Multi-Point Functions

Enote: should directly adapt to multi-point function case

We begin by defining a slightly generalized notion of distributed point functions (DPFs), which accounts for the extra parameter $\mathbb{G}'$. Yaxin: What is $\mathbb{G}'$?

DEFINITION 1 (DPF [4, 9]). *A (2-party) distributed point function (DPF) is a triple of algorithms* $\Pi = (\mathsf{Gen}, \mathsf{Eval}_0, \mathsf{Eval}_1)$ *with the following syntax:*

- $\mathsf{Gen}(1^\lambda, \hat{f}_{\alpha,\beta}) \to (k_0, k_1)$: *On input security parameter $\lambda \in \mathbb{N}$ and point function description $\hat{f}_{\alpha,\beta} = (N, \hat{\mathbb{G}}, \alpha, \beta)$, the (randomized) key generation algorithm* Gen *returns a pair of keys $k_0, k_1 \in \{0, 1\}^*$.* *Yaxin: Matan points out: we want efficient*

procedures, i.e., $|k_b| \in \mathrm{poly}(\lambda)$. Stress it here or add efficiency requirement?We assume that N and $\mathbb{G}$ are determined by each key.

- $\mathsf{Eval}_b(k_b, x) \to y_b$: *On input key $k_b \in \{0, 1\}^*$ and input $x \in [N]$ the (deterministic) evaluation algorithm of server b,* Eval_b *returns $y_b \in \mathbb{G}$.*

We require Π to satisfy the following requirements:

- **Correctness:** *For every λ, $\hat{f} = \hat{f}_{\alpha,\beta} = (N, \hat{\mathbb{G}}, \alpha, \beta)$ such that $\beta \in \mathbb{G}$, and $x \in [N]$, for $b = 0, 1$,*

$$\Pr\left[(k_0, k_1) \leftarrow \mathsf{Gen}(1^\lambda, \hat{f}), \sum_{i=0}^{1} \mathsf{Eval}_b(k_b, x) = f_{\alpha,\beta}(x)\right] = 1$$

- **Security:** *Consider the following semantic security challenge experiment for corrupted server $b \in \{0, 1\}$:*
  (1) *The adversary produces two point function descriptions ($\hat{f}^0 = (N, \hat{\mathbb{G}}, \alpha_0, \beta_0), \hat{f}^1 = (N, \hat{\mathbb{G}}, \alpha_1, \beta_1)) \leftarrow \mathcal{A}(1^\lambda)$, where $\alpha_b \in [N]$ and $\beta_b \in \mathbb{G}$.*
  (2) *The challenger samples $b \leftarrow \{0, 1\}$ and $(k_0, k_1) \leftarrow \mathsf{Gen}(1^\lambda, \hat{f}^b)$.*
  (3) *The adversary outputs a guess $b' \leftarrow \mathcal{A}(k_b)$.*
  Denote by $\mathsf{Adv}(1^\lambda, \mathcal{A}, i) = \Pr[b = b'] - 1/2$ the advantage of $\mathcal{A}$ in guessing b in the above experiment. For every non-uniform polynomial time adversary $\mathcal{A}$ there exists a negligible function ν such that $\mathsf{Adv}(1^\lambda, \mathcal{A}, i) \le \nu(\lambda)$ for all $\lambda \in \mathbb{N}$.

DEFINITION 2 (DMPF). *A (2-party) distributed multi-point function (DMPF) is a triple of algorithms* $\Pi = (\mathsf{Gen}, \mathsf{Eval}_0, \mathsf{Eval}_1)$ *with the following syntax:*

- $\mathsf{Gen}(1^\lambda, \hat{f}_{A,B}) \to (k_0, k_1)$: *On input security parameter $\lambda \in \mathbb{N}$ and point function description $\hat{f}_{A,B} = (N, \hat{\mathbb{G}}, t, A, B)$, the (randomized) key generation algorithm* Gen *returns a pair of keys $k_0, k_1 \in \{0, 1\}^*$.* *Yaxin: On Matan's behalf: same comment as well. Maybe $|k_i| = \mathrm{poly}(\lambda, t)$.*
- $\mathsf{Eval}_b(1^\lambda, k_b, x) \to y_b$: *On input key $k_b \in \{0, 1\}^*$ and input $x \in [N]$ the (deterministic) evaluation algorithm of server b,* Eval_b *returns $y_b \in \mathbb{G}$.*

We require Π to satisfy the following requirements:

- **Correctness:** *For every λ, $\hat{f} = \hat{f}_{A,B} = (N, \hat{\mathbb{G}}, t, A, B)$ such that $B \in \mathbb{G}^t$, and $x \in [N]$, for $b = 0, 1$,*

$$\Pr\left[(k_0, k_1) \leftarrow \mathsf{Gen}(1^\lambda, \hat{f}), \sum_{i=0}^{1} \mathsf{Eval}_b(k_b, x) = f_{A,B}(x)\right] = 1$$

- **Security:** *Consider the following semantic security challenge experiment for corrupted server $b \in \{0, 1\}$:*
  (1) *The adversary produces two t-point function descriptions ($\hat{f}^0 = (N, \hat{\mathbb{G}}, t, A_0, B_0), \hat{f}^1 = (N, \hat{\mathbb{G}}, t, A_1, B_1)) \leftarrow \mathcal{A}(1^\lambda)$, where $\alpha_b \in [N]$ and $\beta_b \in \mathbb{G}$.*
  (2) *The challenger samples $b \leftarrow \{0, 1\}$ and $(k_0, k_1) \leftarrow \mathsf{Gen}(1^\lambda, \hat{f}^b)$.*
  (3) *The adversary outputs a guess $b' \leftarrow \mathcal{A}(k_b)$.*

Denote by $\mathsf{Adv}(1^\lambda, \mathcal{A}, i) = \Pr[b = b'] - 1/2$ the advantage of $\mathcal{A}$ in guessing b in the above experiment. For every non-uniform polynomial time adversary $\mathcal{A}$ there exists a negligible function ν such that $\mathsf{Adv}(1^\lambda, \mathcal{A}, i) \le \nu(\lambda)$ for all $\lambda \in \mathbb{N}$.

We will also be interested in applying the evaluation algorithm on *all* inputs. Given a DMPF $(\mathsf{Gen}, \mathsf{Eval}_0, \mathsf{Eval}_1)$, we denote by $\mathsf{FullEval}_b$ an algorithm which computes Eval_b on every input x. Hence, $\mathsf{FullEval}_b$ receives only a key k_b as input.

2.3 Batch Code

We introduce batch code and probabilistic batch code, which can be used to construct DMPF (see construction 2).

DEFINITION 3 (BATCH CODE[10]). *An (N, M, t, m)-batch code over alphabet Σ is given by a pair of efficient algorithms $(\mathsf{Encode}, \mathsf{Decode})$ such that:*

- $\mathsf{Encode}(x \in \Sigma^N) \to (C_1, C_2, \cdots, C_m)$: *Any string $x \in \Sigma^N$ is encoded into an m-tuple of strings $C_1, C_2, \cdots C_m \in \Sigma^*$ (called buckets) of total length M.*
- $\mathsf{Decode}(I, C_1, C_2, \cdots, C_m) \to \{x[i]\}_{i \in I}$: *On input a set I of t distinct indices in $[N]$ and m buckets, recover t coordinates of x indexed by I by reading at most one coordinate from each of the m buckets.*

An (N, M, t, m)-batch code can be represented by an (N, m)-bipartite graph $G = (U, V, E)$ where each edge $(u_i, v_j) \in E$ corresponds to Encode assigning $x[i]$ to the bucket C_j, while it is guaranteed that any subset $S \subseteq U$ such that $|S| = t$ has a perfect matching to V. Yaxin: Add example instantiation (random regular bipartite graph) and explain it is not efficient.

DEFINITION 4 (PROBABILISTIC BATCH CODE (PBC)[1]). *An (N, M, t, m, ϵ)-probabilistic batch code over alphabet Σ is a randomized (N, M, t, m)-batch code with **public randomness** r (and possibly private randomness for each sub-procedure) such that for any string x and any set I of t distinct indices in $[N]$,*

$$\Pr[\mathsf{Decode}_r(I, \mathsf{Encode}_r(x)) \to \{x[i]\}_{i \in I}] = 1 - \epsilon$$

where the probability is taken over the public randomness r and private randomness of Encode and Decode algorithms.

We mention Cuckoo hashing algorithm[12] as a concrete instantiation of PBC[1].

w-way cuckoo hashing. Given t balls, $m = et$ buckets (e is some expansion parameter that is bigger than 1), and w independent hash functions $h_1, h_2, \cdots, h_w$ randomly mapping every ball to a bucket, allocates all balls to the buckets such that each bucket contains at most one ball through the following process:

1. Choose an arbitrary unallocated ball b. If there is no unallocated ball, output the allocation.
2. Choose a random hash function h_i compute the bucket index $h_i(b)$. If this bucket is empty, then allocate b to this bucket and go to step 1. If this bucket is not empty and filled with ball b', then evict b', allocate b to this bucket set b' the current unallocated ball, and repeat step 2.

If the algorithm terminates then its output is an allocation of balls to buckets such that each bucket contains at most one ball. However

there is no guarantee that the algorithm will terminate - it may end up in a loop and keeps running forever. To fixed this problem, the algorithm should be given a fixed amount of time to run, or equipped with a loop detection process to guarantee termination. We call it a *failure* whenver the algorithm fails to output a proper allocation where each bucket contains at most one ball.

Yaxin: Add asymptotic parameters? Also find evident theorem saying cuckoo hashing is efficient.

The failure probability of cuckoo hashing. Let's denote the failure probability of w-way cuckoo hashing to be $\epsilon = 2^{-\lambda_{\text{stat}}}$. In practice we usually consider the statistical security parameter λ_{stat} to be 30 or 40. The empirical result in [5] shows for $w = 3$, $m = 16384$, $\lambda_{\text{stat}} = 124.4e - 144.6$ where e is the expansion parameter that $m = et$. For $w = 3$, $m = 8192$, $\lambda_{\text{stat}} = 125e - 145$. However we use cuckoo hashing to construct DMPF for t-point functions, in which case we'd also care about t being small, say 2, 3 or 100, and m should not be too large. In this sense the previous empirical results are not complete. Yaxin: [1] uses $w = 3$, $e = 1.5$, $t > 200$ and $\lambda_{\text{stat}} \approx 40$ and claims it follows the analysis from [5], but I don't see how...

The balls, buckets and hash functions can be represented by a w-regular (t, m)-bipartite graph $G = (U, V, E)$ where each left node has w neighbors, and each edge $(u_i, v_j) \in E$ corresponds to $h_l(i) = j$ for some $1 \le l \le w$. In this graph representation the w-way cuckoo hashing essentially computes a perfect matching from U to V. Therefore one can construct a PBC from cuckoo hashing.

CONSTRUCTION 1 (PBC FROM CUCKOO HASHING). *Given w-way cuckoo hashing as a sub-procedure allocating t balls to m buckets with failure probability ϵ, an (N, wN, t, m, ϵ)-PBC is as follows:*

- $\mathsf{Encode}_r(x \in \Sigma^N) \to (C_1, \cdots, C_m)$: *Use r to determine w independent random hash functions $h_1, h_2, \cdots h_w$ that maps from $[N]$ to $[m]$. Initialize $C_1, \cdots, C_m$ to be empty. Let C_j contain $\{x[i] : h_l(i) = j$ for some $l \in [w]\}$, in ascending order of i.*
- $\mathsf{Decode}_r(I, C_1, \cdots, C_m) \to \{x[i]\}_{i \in I}$: *Determine $h_1, \cdots, h_w$ as in Encode. For I of size t, allocate I to $[m]$ using w-way cuckoo hashing. For each $i \in I$, fetch $x[i]$ from the C_j where i is allocated to C_j. Note that $x[i]$ can be found in the kth position of C_j where i is the kth smallest index of $\{i : h_l(i) = j$ for some $l \in [w]\}$.*

2.4 DMPF Construction from (Probabilistic) Batch Code

We display the construction of DMPF from black-box usage of DPF basing on PBC with appropriate parameters, which has been discussed in previous literature[2, 6].

CONSTRUCTION 2 (DMPF FROM DPF). *Given DPF for any domain of size no larger than N and output group $\mathbb{G}$, and an (N, M, t, m, ϵ)-PBC with alphabet $\Sigma = \mathbb{G}$, we can construct a DMPF scheme for t-point functions with domain size N and output group $\mathbb{G}$ as follows:*

- $\mathsf{Gen}(1^\lambda, \hat{f}_{A,B}) \to (k_0, k_1)$: *Suppose $A = \{\alpha_1, \cdots, \alpha_t\}$ and $B = \{\beta_1, \cdots, \beta_t\}$. Let $TT \in \mathbb{G}^N$ be the truth table of $\hat{f}_{A,B}$. Compute $\mathsf{Encode}(TT) \to (C_1, \cdots, C_m)$ according to the PBC. Then run $\mathsf{Decode}(A, C_1, \cdots, C_m)$ to determine a perfect matching from*

A to $\{C_1, \cdots, C_m\}$. For $1 \le i \le m$, let $f_i : [|C_i|] \to \mathbb{G}$ be the following:

- If C_i is assigned none of A by the perfect matching, then set f_i to be the all-zero function.
- If exactly one α_j of A is assigned to the lth position of C_i, then set f_i to be the point function that outputs β_j on l and 0 elsewhere.

For $1 \le i \le m$, invoke $DPF.\text{Gen}(1^\lambda, f_i) \to (k_0^i, k_1^i)$. Set $(k_0, k_1) = (\{k_0^i\}_{i \in [m]}, \{k_1^i\}_{i \in [m]})$. If Decode fails then run Encode and Decode again with fresh randomness.

- $\text{Eval}_b(k_b, x) \to y_b$: Follow $\text{Encode}(TT)$ to determine the positions $l_{j_1}, l_{j_2}, \cdots, l_{j_s}$ such that the xth entry of TT is sent to the l_{j_i}-th position of C_{j_i}. Compute $y_b = \sum_{i=1}^{s} DPF.\text{Eval}_b(k_b^{j_i}, l_i)$.

The scheme is correct with overwhelming probability and has distinguish advantage $< 2\epsilon$.

Note that if one use batch code instead of PBC then the DMPF scheme perfectly correct and secure. When instantiating PBC from w-way cuckoo hashing, the *key generation time* is roughly the time needed for computing cuckoo hashing algorithm plus the total time of all $DPF.\text{Gen}(1^\lambda, f_i)$. The *evaluation time* is roughly the total time of all $DPF.\text{Eval}_b(k_b^{j_i}, l_i)$. Similarly, the *full-domain evaluation time* is roughly the total time of all $DPF.\text{FullEval}_b(k_b^j)$ for $j = 1, \ldots, m$.

2.5 Oblivious Key-Value Stores

We introduce the notion of Oblivious key-value stores (OKVS) which can be used to construct DMPF. OKVS was originally proposed as a primitive for private set intersection (PSI) protocols (see [8, 13]).

DEFINITION 5 (OBLIVIOUS KEY-VALUE STORES (OKVS)[8, 13]). *An Oblivious Key-Value Stores scheme is a pair of randomized algorithms* $(\text{Encode}_r, \text{Decode}_r)$ *with respect to a statistical security parameter* λ_{stat} *and a computational security parameter* λ, *a randomness space* $\{0,1\}^\kappa$, *a key space* $\mathcal{K}$, *a value space* $\mathcal{V}$, *input length* t *and output length* m. *The algorithms are of the following syntax:*

- *$\text{Encode}_r(\{(k_1, v_1), (k_2, v_2), \cdots, (k_t, v_t)\}) \to P$: On input t key-value pairs with distinct keys, the encode algorithm with randomness r in the randomness space outputs an encoding $P \in \mathcal{V}^m \cup \bot$.*
- *$\text{Decode}_r(P, k) \to v$: On input an encoding from $\mathcal{V}^m$ and a key $k \in \mathcal{K}$, output a value v.*

We require the scheme to satisfy

- *For all $S \in (\mathcal{K} \times \mathcal{V})^t$, $\Pr_{r \leftarrow \{0,1\}^\kappa}[\text{Encode}_r(S) = \bot] \le 2^{-\lambda_{\text{stat}}}$.*
- *For all $S \in (\mathcal{K} \times \mathcal{V})^t$ and $r \in \{0,1\}^\kappa$ such that $\text{Encode}_r(S) \to P \ne \bot$, it is the case that $\text{Decode}_r(P, k) \to v$ whenever $(k, v) \in S$.*
- ***Obliviousness:** Given any distinct key sets $\{k_1^0, k_2^0, \cdots, k_t^0\}$ and $\{k_1^1, k_2^1, \cdots, k_t^1\}$ that are different, if they are paired with random values then their encodings are computationally indistinguishable, i.e.,*

$$\{r, \text{Encode}_r(\{(k_1^0, v_1), \cdots, (k_t^0, v_t)\})\}_{v_1, \cdots, v_t \leftarrow \mathcal{V}, r \leftarrow \{0,1\}^\kappa}$$

$$\approx_c \{r, \text{Encode}_r(\{(k_1^1, v_1), \cdots, (k_t^1, v_t)\})\}_{v_1, \cdots, v_t \leftarrow \mathcal{V}, r \leftarrow \{0,1\}^\kappa}$$

One can obtain a linear OKVS *if in addition require:*

- ***Linearity:** There exists a function family $\{\text{row}_r : \mathcal{K} \to \mathcal{V}^m\}_{r \in \{0,1\}^\kappa}$ such that $\text{Decode}_r(P, k) = \langle \text{row}_r(k), P \rangle$.*

The Encode process for a linear OKVS is the process of sampling a random P from the set of solutions of the linear system $\{\langle \text{row}_r(k_i), P \rangle = v_i\}_{1 \le i \le t}$.

We evaluate an OKVS scheme by its encoding size (output length m), encoding time and decoding time. We stress the following two (linear) OKVS constructions:

CONSTRUCTION 3 (POLYNOMIAL). *Suppose $\mathcal{K} = \mathcal{V} = \mathbb{F}$ is a field. Set*

- *$\text{Encode}(\{(k_i, v_i)\}_{1 \le i \le t}) \to P$ where P is the coeffients of a $(t-1)$-degree $\mathbb{F}$-polynomial g_P that $g_P(k_i) = v_i$ for $1 \le i \le t$.*
- *$\text{Decode}(P, k) \to g_P(k)$.*

The polynomial OKVS possesses an optimal encoding size $m = n$, but the Encode process is a polynomial interpolation which is only known to be achieved in time $O(t \log^2 t)$. The time for a single decoding is $O(t)$ and that for batched decodings is (amortized) $O(\log^2 t)$.

An alternative construction that has near optimal encoding size but much better running time is as follows.

CONSTRUCTION 4 (3-HASH GARBLED CUCKOO TABLE (3H-GCT)[8, 13]). *Suppose $\mathcal{V} = \mathbb{F}$ is a field. Set $\text{row}_r(k) := \text{row}_r^{\text{sparse}}(k) || \text{row}_r^{\text{dense}}(k)$ where $\text{row}_r^{\text{sparse}}(k)$ outputs a uniformly random weight-w vector in $\{0,1\}^{m_1}$, and $\text{row}_r^{\text{dense}}(k)$ outputs a short dense vector in $\mathbb{F}^{m_2}$.*

- *$\text{Encode}(\{(k_i, v_i)\}_{1 \le i \le t}) \to P$ where P is solved from the system $\{\langle \text{row}_r(k_i), P \rangle = v_i\}_{1 \le i \le t}$ using the triangulation algorithm in [13].*
- *$\text{Decode}(P, k) \to \langle \text{row}_r(k), P \rangle$.*

We denote $m_1 = et$, where e is an expandion parameter indicating the rough blowup to store t pairs. In practice the number of dense columns m_2 is usually set to a small constant.

This OKVS construction features a linear encoding time, constant decoding time (the constant relates to w and m_2) while having a linear encoding size.

Yaxin: TBD: Carefully(!) recompute the comparison table for OKVS.

We'll mostly use the expansion parameter e and the number of dense columns $m_2 := \hat{g}$ (where $\hat{g}$ is a parameter relating to the equation system solving process) according to the analysis in [13]: Given w, t and λ_{stat}, the choices of the e and $\hat{g}$ are fixed through the following steps:

- Set $e^* = \begin{cases} 1.223 & w = 3 \\ 1.293 & w = 4 \\ 0.1485w + 0.6845 & w \ge 5 \end{cases}$.
- Compute $\alpha := 0.55 \log_2 t + 0.093w^3 - 1.01w^2 + 2.92w - 0.13$.
- $e := e^* + 2^{-\alpha}(\lambda_{\text{stat}} + 9.2)$.
- $\hat{g} := \frac{\lambda_{\text{stat}}}{(w-2)\log_2(et)}$.

Yaxin: Fix t and λ_{stat}, we want to find the best choice of w. The adavantageous choices of w in [13] are $w = 3$ and $w = 5$. From the first sight when w is smaller e can be smaller but $\hat{g}$ will be larger. Since $w + \hat{g}$ stands for number of $\mathbb{F}$-ADD's and $\hat{g}$ stands for number of $\mathbb{F}$-MULT's in decoding, previously I thought $\hat{g}$ is the dominating

factor of Decode running time. However table 1 in [13] suggests that $w = 3$ outruns nearly all of other choices of w while $w = 5$ is almost 3 times slower in decoding time. This may suggest there are some other heavy computations other than $\mathbb{F}$-MULT that need to be considered when evaluating running time.

The range of t previous literature [8, 13] have considered in their empirical results are also limited, which will be one of our problems. We want to cover small t, say $t < 100$, while previous literature aiming for constructing PSI protocols usually consider very large t.

One may also let row_r^{dense} output a short dense vector in $\{0, 1\}^{m_2}$, which avoids multiplication of large field elements in the encoding and decoding processes. To achieve same level of security one could simply set $m_2 = \hat{g} + \lambda_{\text{stat}}$, as proposed in [8, 13]. Yaxin: TBD: mention some connections to cuckoo hashing?

3 NEW DMPF CONSTRUCTIONS

In this section, we display two new constructions of DMPF that follow the same paradigm shown in fig. 1.

We begin by introducing the DMPF paradigm in fig. 1, which is based on the idea of the DPF construction in [4]. Each key $k_b (b = 0, 1)$ generated by $\text{Gen}(\hat{f}_{A,B})$ can span a height-n (n is the input length of $\hat{f}_{A,B}$) complete binary tree T_b (call it the evaluation tree) that has 2^n leaf nodes, and each of its nodes is associated with a value. Party b can evaluate the input $x = x_1 \cdots x_n$ by starting from the root of the evaluation tree, on the ith layer going left if $x_i = 0$ and going right if $x_i = 1$, until reaching a leaf node then computing the result according to the value on this leaf node.

Each node of this tree is associated with a λ-bit seed string and a l-bit sign string. For a parent node on the ith layer with seed and sign, its children's seeds and signs are generated by $\text{PRG}(\text{seed}) \oplus$ Correction, where the Correction is determined by the parent node's position, its sign and a correction word $CW^{(i)}$ associated with that layer. On a leaf node on the last layer, its seed will generate a random element in the output group, which will be corrected by adding a Correction determined by the leaf node's position, its sign and the last correction word $CW^{(n+1)}$.

The computation of values on evaluation trees T_0, T_1 involves the following methods in the paradigm in fig. 1:

- Initialize computes the values on the roots of T_0, T_1.
- GenCW computes correction words associated with the first n layers, namely $\{CW^{(1)}, \cdots CW^{(n)}\}$. Two parties use the same set of correction words.
- GenConvCW computes the correction word $CW^{(n+1)}$ associated with the last layer that'll help genereate the final result in the output group $\mathbb{G}$. Two parties use the same set of correction words.
- Correct on input a node's position, its sign and the correction word $CW^{(i)}$ corresponding to its layer, outputs an additive correction for its children's value.
- ConvCorrect on input a leaf node's position, its sign and the correction word $CW^{(n+1)}$, outputs an additive correction for the final result in the output group $\mathbb{G}$.

Call any path from the root to a leaf (corresponding to a string in A) an accepting path. To force the correctness, we maintain the following invariance on the evaluation trees T_0, T_1 of the two parties:

- If a node is not on any accepting path, then in T_0 and T_1 it is associated with the same seed and sign.
- If a node is on an accepting path, then T_0 and T_1 assign to it with different signs that controls the corrections on its children (or on the output if it's a leaf node).

If a node is not on any accepting path, then its children are associated with the same values in T_0 and T_1, and so is the subtree rooted at this node. Hence one only needs to design how to correct children of nodes on the accepting paths. We provides the big-state DMPF and the OKVS-based DMPF basing on this paradigm, that differ mainly on the way of generating and computing corrections.

Figure 1: The paradigm of our DMPF schemes. We leave the sign string length l, methods Initialize, GenCW, GenConvCW, Correct, ConvCorrect **to be determined by specific constructions. In the paradigm** $A^{(i)}$ **denotes the list of accepting paths reaching the ith layer, and** seed$^{(i)}$ and sign$^{(i)}$ **denote the lists of** seed **values and** sign **values at the nodes on the ith layer that are on accepting paths.**

Public parameters:
The t-point function family $\{f_{A,B}\}$ with t an upperbound of the number of nonzero points, input domain $[N] = \{0,1\}^n$ and the output group $\mathbb{G}$.
Suppose there is a public PRG $G : \{0,1\}^\lambda \to \{0,1\}^{2\lambda+2l}$. Parse $G(x) = G_0(x)\|G_1(x)$ to the left half and right half of the output.
Suppose there is a public PRG $G_{\text{convert}} : \{0,1\}^\lambda \to \mathbb{G}$.

procedure GEN$(1^\lambda, \hat{f}_{A,B})$
    Denote $A = (\alpha_1, \cdots, \alpha_t)$ in lexicographically order, $B = (\beta_1, \cdots, \beta_t)$. If $|A| < t$, extend A to size-t with arbitrary $\{0,1\}^n$ strings and B with 0's.
    For $0 \le i \le n-1$, let $A^{(i)}$ denote the sorted and deduplicated list of i-bit prefixes of strings in A. Specifically, $A^{(0)} = [\epsilon]$.
    For $0 \le i \le n-1$ and $b = 0,1$, initialize empty lists seed$_b^{(i)}$ and sign$_b^{(i)}$.
    Initialize($\{$seed$_b^{(0)}$, sign$_b^{(0)}\}_{b=0,1}$).
    for $i = 1$ to n **do**
        $CW^{(i)} \leftarrow$ GenCW$(i, A, \{$seed$_b^{(i-1)}$, sign$_b^{(i-1)}\}_{b=0,1})$.
        for $k = 1$ to $|A^{(i-1)}|$ and $z = 0,1$ **do**
            Compute $C_{\text{seed},b}\|C_{\text{sign}^0,b}\|C_{\text{sign}^1,b} \leftarrow$ Correct$(A^{(i-1)}[k], \text{sign}_b^{(i-1)}[k], CW^{(i)})$ for $b = 0,1$.
            if $A^{(i-1)}[k]\|z \in A^{(i)}$ **then**
                Append the first λ bit of $G_z(\text{seed}_b^{(i-1)}[k]) \oplus (C_{\text{seed},b}\|C_{\text{sign}^z,b})$ to seed$_b^{(i)}$ and the rest to sign$_b^{(i)}$.
            end if
        end for
    end for
    $CW^{(n+1)} \leftarrow$ GenConvCW$(A, B, \{$seed$_b^{(n)}$, sign$_b^{(n)}\}_{b=0,1})$.
    Set $k_b \leftarrow (\text{seed}_b^{(0)}, \text{sign}_b^{(0)}, CW^{(1)}, CW^{(2)}, \cdots, CW^{(n+1)})$.
    return (k_0, k_1).
end procedure

procedure EVAL$_b(1^\lambda, k_b, x)$
    Parse $k_b = ([\text{seed}], [\text{sign}], CW^{(1)}, CW^{(2)}, \cdots, CW^{(n+1)})$.
    Denote $x = x_1 x_2 \cdots x_n$.
    for $i = 1$ to n **do**
        $C_{\text{seed}}\|C_{\text{sign}^0}\|C_{\text{sign}^1} \leftarrow$ Correct$(x_1 \cdots x_{i-1}, \text{sign}, CW^{(i)})$.
        seed$\|$sign $\leftarrow G_{x_i}(\text{seed})$.
        seed$\|$sign $\leftarrow G_{x_i}(\text{seed}) \oplus (C_{\text{seed}}\|C_{\text{sign}^{x_i}})$.
    end for
    return $(-1)^b \cdot \big(G_{\text{convert}}(\text{seed}) + \text{ConvCorrect}(x, \text{sign}, CW^{(n+1)})\big)$.
end procedure

procedure FULLEVAL$_b(1^\lambda, k_b)$
    Parse $k_b = (\text{seed}^{(0)}, \text{sign}^{(0)}, CW^{(1)}, CW^{(2)}, \cdots, CW^{(n+1)})$.
    For $1 \le i \le n$, Path$^{(i)} \leftarrow$ the lexicographical ordered list of $\{0,1\}^i$. Path$^{(0)} \leftarrow [\epsilon]$.
    Yaxin: The evaluation is BFS-style, which costs a lot of memory to store lists seed$^{(i)}$, sign$^{(i)}$. Need a DFS version for large N to reduce memory use? Write in the clear or explain by words?
    for $i = 1$ to n **do**
        for $k = 1$ to 2^{i-1} **do**
            $C_{\text{seed}}\|C_{\text{sign}^0}\|C_{\text{sign}^1} \leftarrow$ Correct$(\text{Path}^{(i-1)}[k], \text{sign}^{(i-1)}[k], CW^{(i)})$.
            seed$^{(i)}[2k]\|$sign$^{(i)}[2k] \leftarrow G_0(\text{seed}^{(i-1)}[k]) \oplus (C_{\text{seed}}\|C_{\text{sign}^0})$.
            seed$^{(i)}[2k+1]\|$sign$^{(i)}[2k+1] \leftarrow G_1(\text{seed}^{(i-1)}[k]) \oplus (C_{\text{seed}}\|C_{\text{sign}^1})$.
        end for
    end for
    for $k = 1$ to 2^n **do**
        Output$[k] \leftarrow (-1)^b \cdot \big(G_{\text{convert}}(\text{seed}^{(n)}[k]) + \text{ConvCorrect}(\text{Path}[k], \text{sign}^{(n)}[k], CW^{(n+1)})\big)$.
    end for
    return Output.
end procedure

3.1 Big-State DMPF

Displayed in fig. 2. TBD: explain

3.2 OKVS-based DMPF

Displayed in fig. 3. TBD: explain

Figure 2: The parameter l and methods' setting that turns the paradigm of DMPF in fig. 1 into the big-state DMPF.

Set $l \leftarrow t$, the upperbound of $|A|$.
procedure INITIALIZE($\{\text{seed}_b^{(0)}, \text{sign}_b^{(0)}\}_{b=0,1}$)

For $b = 0, 1$, let $\text{seed}_b^{(0)} = [r_b]$ where $r_b \xleftarrow{\$} \{0, 1\}^\lambda$.

For $b = 0, 1$, set $\text{sign}_b^{(0)} = [b \| 0^{t-1}]$.
end procedure

procedure GENCW($i, A, \{\text{seed}_b^{(i-1)}, \text{sign}_b^{(i-1)}\}_{b=0,1}$)

Let $\{A^{(i)}\}_{0 \leq i \leq n}$ be defined as in fig. 1.

Sample a list CW of t random strings from $\{0, 1\}^{\lambda+2t}$.

for $k = 1$ to $|A^{(i-1)}|$ **do**

Parse $G(\text{seed}_b^{(i-1)}[k]) = \text{seed}_b^0 \| \text{sign}_b^0 \| \text{seed}_b^1 \| \text{sign}_b^1$, for $b = 0, 1$, $\text{seed}_b^0, \text{seed}_b^1 \in \{0, 1\}^\lambda$ and $\text{sign}_b^0, \text{sign}_b^1 \in \{0, 1\}^t$.

Compute $\Delta\text{seed}^c = \text{seed}_0^c \oplus \text{seed}_1^c$ and $\Delta\text{sign}^c = \text{sign}_0^c \oplus \text{sign}_1^c$ for $c = 0, 1$.

Denote $\text{path} \leftarrow A^{(i-1)}[k]$.

if both $\text{path} \| z$ for $z = 0, 1$ are in $A^{(i)}$ **then**

$d \leftarrow$ the index of $\text{path} \| 0$ in $A^{(i)}$.

$CW[k] \leftarrow r \| \Delta\text{sign}^0 \oplus e_d \| \Delta\text{sign}^1 \oplus e_{d+1}$ where $r \xleftarrow{\$} \{0, 1\}^\lambda$, $e_d = 0^{d-1} 1 0^{t-d}$.

else

Let z be such that $\text{path} \| z \in A^{(i)}$.

$d \leftarrow$ the index of $\text{path} \| z$ in $A^{(i)}$.

$CW[k] \leftarrow \begin{cases} \Delta\text{seed}^1 \| \Delta\text{sign}^0 \oplus e_d \| \Delta\text{sign}^1 & z = 0 \\ \Delta\text{seed}^0 \| \Delta\text{sign}^0 \| \Delta\text{sign}^1 \oplus e_d & z = 1 \end{cases}$.

end if

end for

return CW.
end procedure

procedure GENCONVCW($A, B, \{\text{seed}_b^{(n)}, \text{sign}_b^{(n)}\}$)

Sample a list CW of t random $\mathbb{G}$-elements.

for $k = 1$ to $|A|$ **do**

$\Delta g \leftarrow G_{\text{convert}}(\text{seed}_0^{(n)}[k]) - G_{\text{convert}}(\text{seed}_1^{(n)}[k])$.

$CW[k] \leftarrow (-1)^{\text{sign}_0^{(n)}[k][k]}(\Delta g - B[k])$.

end for

return CW.
end procedure

procedure CORRECT($\bar{x}, \text{sign}, CW$)

return $C_{\text{seed}} \| C_{\text{sign}^0} \| C_{\text{sign}^1} \leftarrow \sum_{i=1}^t \text{sign}[i] \cdot CW[i]$, where C_{sign^0} and C_{sign^1} are t-bit.
end procedure

procedure CONVCORRECT(x, sign, CW)

return $\sum_{i=1}^t \text{sign}[i] \cdot CW[i]$.
end procedure

Figure 3: The parameter l and methods' setting that turns the paradigm of DMPF in fig. 1 into the OKVS-based DMPF.

Set $l \leftarrow 1$.

For $1 \leq i \leq n$, let OKVS_i be an OKVS scheme (definition 5) with key space $\mathcal{K} = \{0,1\}^{i-1}$, value space $\mathcal{V} = \{0,1\}^{\lambda+2}$ and input length t.

let $\mathrm{OKVS}_{\mathrm{convert}}$ be an OKVS scheme with key space $\mathcal{K} = \{0,1\}^n$, value space $\mathcal{V} = \mathbb{G}$ and input length t.

procedure INITIALIZE($\{\mathrm{seed}_b^{(0)}, \mathrm{sign}_b^{(0)}\}_{b=0,1}$)

    For $b = 0, 1$, let $\mathrm{seed}_b^{(0)} = [r_b \xleftarrow{\$} \{0,1\}^\lambda]$ and $\mathrm{sign}_b^{(0)} = [b]$.

end procedure

procedure GENCW($i, A, \{\mathrm{seed}_b^{(i-1)}, \mathrm{sign}_b^{(i-1)}\}_{b=0,1}$)

    Let $\{A^{(i)}\}_{0 \leq i \leq n}$ be defined as in fig. 1.

    Sample a list V of t random strings from $\{0,1\}^{\lambda+2}$.

    for $k = 1$ to $|A^{(i-1)}|$ **do**

        Parse $G(\mathrm{seed}_b^{(i-1)}[k]) = \mathrm{seed}_b^0 \| \mathrm{sign}_b^0 \| \mathrm{seed}_b^1 \| \mathrm{sign}_b^1$, for $b = 0, 1$, $\mathrm{seed}_b^0, \mathrm{seed}_b^1 \in \{0,1\}^\lambda$ and $\mathrm{sign}_b^0, \mathrm{sign}_b^1 \in \{0,1\}$.

        Compute $\Delta\mathrm{seed}^c = \mathrm{seed}_0^c \oplus \mathrm{seed}_1^c$ and $\Delta\mathrm{sign}^c = \mathrm{sign}_0^c \oplus \mathrm{sign}_1^c$ for $c = 0, 1$.

        Denote path $\leftarrow A^{(i-1)}[k]$.

        if both path$\|z$ for $z = 0, 1$ are in $A^{(i)}$ **then**

            $V[k] \leftarrow r \| \Delta\mathrm{sign}^0 \oplus 1 \| \Delta\mathrm{sign}^1 \oplus 1$, where $r \xleftarrow{\$} \{0,1\}^\lambda$.

        else

            Let z be such that path$\|z \in A^{(i)}$.

            $V[k] \leftarrow \Delta\mathrm{seed}^1 \| \Delta\mathrm{sign}^0 \oplus (1 - z) \| \Delta\mathrm{sign}^1 \oplus z$.

        end if

    end for

    return $\mathrm{OKVS}_i.\mathrm{Encode}(\{A^{(i-1)}[k], V[k]\}_{1 \leq k \leq |A^{(i-1)}|})$.

end procedure

procedure GENCONVCW($A, B, \{\mathrm{seed}_b^{(n)}, \mathrm{sign}_b^{(n)}\}$)

    Sample a list V of t random $\mathbb{G}$-elements.

    for $k = 1$ to $|A|$ **do**

        $\Delta g \leftarrow G_{\mathrm{convert}}(\mathrm{seed}_0^{(n)}[k]) - G_{\mathrm{convert}}(\mathrm{seed}_1^{(n)}[k])$.

        $V[k] \leftarrow (-1)^{\mathrm{sign}_0^{(n)}[k][k]}(\Delta g - B[k])$.

    end for

    return $\mathrm{OKVS}_{\mathrm{convert}}(\{A[k], V[k]\}_{1 \leq k \leq t})$.

end procedure

procedure CORRECT($\bar{x}, \mathrm{sign}, CW$)

    return $C_{\mathrm{seed}} \| C_{\mathrm{sign}^0} \| C_{\mathrm{sign}^1} \leftarrow \mathrm{sign} \cdot \mathrm{OKVS}_i.\mathrm{Decode}(CW, \bar{x})$, where C_{sign^0} and C_{sign^1} are bits.

end procedure

procedure CONVCORRECT(x, sign, CW)

    return $\mathrm{sign} \cdot \mathrm{OKVS}_{\mathrm{convert}}.\mathrm{Decode}(CW, x)$.

end procedure

Yaxin: One point: the row matrix of the current layer contains the row matrix of the previous layers, which might be useful for speedup.

3.3 Comparison

Table 1 displays the keysize, running time of Gen,Eval and FullEval for different DMPF schemes, computed in terms of costs of abstract tools such as PRG, batch code and OKVS. Yaxin: We can plug in the actual costs of these tools after carrying out a complete experiment. Yaxin: Note that the PRG in big-state DMPF maps from $\{0,1\}^\lambda$ to $\{0,1\}^{2\lambda+2t}$.

Yaxin: Take PCG as a potential application. We care about FullEval time which is related to PCG seed expanding time. In this aspect, the batch code DMPF consumes $d \times$PRGs than big-state DMPF and OKVS-based DMPF, while big-state DMPF's FullEval time scales with t and OKVS-based DMPF in addition consumes large field multiplications (in OKVS decoding, and maybe more than this). Therefore we expect different DMPF schemes to be the top choice in different choices of t and depending on the computing time of PRG and large field multiplication.

3.4 Distributed Key Generation

Table 1: Keysize and running time comparison for different DMPF constructions for domain size N, t accepting points, output group $\mathbb{G}$ and computational security parameter λ. We leave this table with the abstraction of (probabilistic) batch code in the second column and the abstraction of OKVS in the last column, and plug in concrete instantiations later. m in the second column stands for the number of buckets used in batch code, and w stands for the number of buckets that each input coordinate is mapped to (we only consider regular degree because this is the case in most instantiations). Yaxin: Denote T_G as the time for computing $G : \{0,1\}^{\lambda+1} \rightarrow \{0,1\}^{2\lambda+2}$, and $T_{G_{convert}}$ as the time for computing $G_{convert} : \{0,1\}^{\lambda} \rightarrow \mathbb{G}$. In the last column, denote OKVS as the OKVS scheme used for the first n correction words, and OKVS_{conv} as the OKVS scheme used for the last correction word.

	Sum of t DPFs	Batch code DMPF[1, 2, 6, 14]	Big-state DMPF	OKVS-based DMPF
Keysize	$t(\lambda+2)\log N + t\log\mathbb{G}$	$m(\lambda+2)\log(wN/m) + m\log\mathbb{G}$	$t(\lambda+2t)\log N + t\log\mathbb{G}$	$\log N \times \text{OKVS.Codesize}$ $+\text{OKVS}_{conv}.\text{Codesize}$
$Gen()$	$2t\log N \times T_G + 2t \times T_{G_{convert}}$	$2m\log(wN/m) \times T_G + 2m \times T_{G_{convert}}$ Finding a matching of t inputs to m buckets	$2t\log N \times T_{G^*}$ [1]	$2t\log N \times T_G + 2t \times T_{G_{convert}}$ $+\log N \times \text{OKVS.Encode}$ $+\text{OKVS}_{conv}.\text{Encode}$
$Eval()$	$t\log N \times T_G + t \times T_{G_{convert}}$	$w\log(wN/m) \times T_G + w \times T_{G_{convert}}$ Finding all positions the input is mapped to	$\log N \times T_{G^*} + T_{G_{convert}}$	$\log N \times T_G$ $+\log N \times \text{OKVS.Decode}$ $+\text{OKVS}_{conv}.\text{Decode}$
$FullEval()$	$tN \times T_G + tN \times T_{G_{convert}}$	$wN \times T_G + wN \times T_{G_{convert}}$ Finding the full domain in buckets	$N \times T_{G^*} + N \times T_{G_{convert}}$	$N \times T_G$ $+N \times \text{OKVS.Decode}$ $+N \times \text{OKVS}_{conv}.\text{Decode}$

[1] The PRG used in big-state DMPF maps from $\{0,1\}^{\lambda}$ to $\{0,1\}^{2\lambda+2t}$ whose computation time should grow with t. We mark this PRG as G^* and its computation time as T_{G^*}.

4 APPLICATIONS

For convenience of discussion DMPF applications, we use $\text{DMPF}_{t,N,\mathbb{G}}$ to denote a DMPF scheme for t-point functions with domain $[N]$ and output group $\mathbb{G}$.

4.1 PCG for OLE from Ring-LPN

We begin by briefly introducing the protocol of PCG for OLE from Ring-LPN assumption, proposed in [3].

The PCG protocol for OLE correlation. The hardness assumption we will make use of is a variant of Ring-LPN, called module-LPN assumption.

DEFINITION 6 (MODULE-LPN). *Let $c \geq 2$ be an integer, $R = \mathbb{Z}_p[X]/F(X)$ for a prime p and a deg-N polynomial $F(X) \in \mathbb{Z}_p[X]$, and $\mathcal{HW}_{R,t}$ be the uniform distribution over weight-t polynomials in R whose degree is less than N and has at most t nonzero coefficients. For $R = R(\lambda)$, $t = t(\lambda)$ and $m = m(\lambda)$, we say that the module-LPN problem R^c-LPN is hard if for every nonuniform polynomial-time probabilistic distinguisher $\mathcal{A}$, it holds that*

$$|\Pr[\mathcal{A}(\{\vec{a}^{(i)}, \langle\vec{a}^{(i)},\vec{s}\rangle+\vec{e}^{(i)}\}_{i\in[m]})]-\Pr[\mathcal{A}(\{\vec{a}^{(i)},\vec{u}^{(i)}\}_{i\in[m]})] \leq \text{negl}(\lambda)$$

where the probabilities are taken over the randomness of $\mathcal{A}$, random samples $\vec{a}^{(1)}, \cdots, \vec{a}^{(m)} \leftarrow R^{c-1}$, $\vec{u}^{(1)}, \cdots, \vec{u}^{(m)} \leftarrow R$, $\vec{s} \leftarrow \mathcal{HW}_{R,t}^{c-1}$, and $\vec{e}^{(1)}, \cdots, \vec{e}^{(m)} \leftarrow \mathcal{HW}_{R,t}$.

When we only consider $m = 1$, each R^c-LPN instance $\langle\vec{a},\vec{s}\rangle + \vec{e}$ can be restated as $\langle\vec{a}',\vec{e}'\rangle$ where $\vec{a}' = 1||\vec{a}$ and $\vec{e}' \leftarrow \mathcal{HW}_{R,t}^c$.

The PCG protocol in [3] generates seed for the OLE correlation $(x_0, x_1, z_0, z_1) \in R^4$ such that $x_0 + x_1 = z_0 \cdot z_1$. The idea is to first set $z_b = \langle\vec{a},\vec{e_b}\rangle$ (an R^c-LPN instance with public $\vec{a}$ and $\vec{e_b} \leftarrow \mathcal{HW}_{R,t}^c$). Basing on the fact that $\langle\vec{a},\vec{e_0}\rangle \cdot \langle\vec{a},\vec{e_1}\rangle = \langle\vec{a}\otimes\vec{a}, \vec{e_0}\otimes\vec{e_1}\rangle$, the next step is to additively share the tensor product $\vec{e_1} \otimes \vec{e_1}$ and each party can compute an additive share of $z_0 \cdot z_1$. Note that the tensor product $\vec{e_0} \otimes \vec{e_1}$ consists of c^2 entries, each being an deg-2N polynomial

with at most t^2 nonzero coefficients. Therefore it can be shared by invoking $\text{DMPF}_{t^2,2N,\mathbb{Z}_p}$ for c^2 times.

One can compute the seed size and expanding time of this PCG protocol as follows:

- The seed size is $ct(\log N + \log p)$ bits for specifying $\vec{e_b}$ plus the $c^2\times$keysize of DMPF.
- The expanding time is c^2N multiplications in R plus $c^2\times$full-domain evaluation time of DMPF.

REMARK 7. *Note that the above PCG protocol generates seed for OLE correlation over ring R. One can immediately convert an OLE correlation over ring R to N OLE correlations over $\mathbb{Z}_p$ if the polynomial $F(X)$ splits into N distinct linear factors modulo p. Therefore we mostly consider reducible F of such form.*

A previous optimization is to substitute $\mathcal{HW}_{R,t}$ with regular weight-t polynomials denoted as regular-$\mathcal{HW}_{R,t}$. Each regular weight-t polynomial e contains exactly one nonzero coefficient e_j in the range of degree $[j \cdot (N/t), (j+1) \cdot (N/t)-1]$ for $j = 0, \cdots t-1$. When multiplying two regular weight-t polynomials e and f, $e_i \cdot f_j$ contributes to a coefficient in the range of degree $[(i+j)\cdot(2N/t), (i+j+2)\cdot(2N/t)-2]$. Therefore the deg-2$N$ polynomial $e \cdot f$ can be shared by invoking $\{\text{DMPF}_{k,2N/t,\mathbb{Z}_p}\}_{k=1,2,\cdots,t-1,t,t-1,\cdots,2,1}$, which cuts down the domain size compared to the original design.

The previous literature uses sum of DPFs to achieve DMPF in either the original design with nonregular noise or the optimized design with regular noise. It indicates using batch code to achieve DMPF as another optimization but not in the clear. We'll analyze the cost of this PCG protocol under the following settings:

(1) with regular noise and each multiplication of sparse polynomials is shared by 2 sets of $\{\text{DMPF}_{k,2N/t,\mathbb{Z}_p}\}_{1\leq k\leq t-1}$ and $\text{DMPF}_{t,2N/t,\mathbb{Z}_p}$;

(2) with nonregular noise and each multiplication of sparse polynomials is shared by $\text{DMPF}_{t^2,2N,\mathbb{Z}_p}$.

Table 2: Concrete applications of DMPF.

Concrete application	Cost in terms of DMPF per correlation/execution	Typical DMPF parameters
PCG for OLE from Ring-LPN	seedsize $\propto$ DMPF.$keysize$ expand time $\propto$ DMPF.$FullEval()$	$t = 5^2, 16^2, 76^2$ $N = 2^{20}$
PSI-WCA	communication $\propto$ DMPF.$keysize$ client computation $\propto$ DMPF.$Gen()$ server computation $\propto$ DMPF.$Eval()$	$t =$ any $N = 2^{128}$

We'll instantiate DMPF in different ways as listed in table 1. The costs of PCG protocols under different settings are listed in table 3.

Yaxin: One caveat: can batch-code / OKVS-based DMPF fit into the regular design, while it requires shares of 1, 2, 3-point functions?

From table 3 we can see that if we allow small blowup of the seed size of PCG, then we can gain much faster seed expansion by using nonregular noise distribution and either batch-code DMPF or big-state DMPF or OKVS-based DMPF.

Comparing the entropy of regular and nonregular noise. Moreover, let's consider using the entropy of the noise distribution as the security parameter (which means that we consider the adversary's attack being randomly guessing the secret). Since the entropy of the nonregular noise distribution $\mathcal{HW}_{R_1, t_1}$ is $\log(\frac{N_1^{t_1}}{t_1!})$ while the entropy of the regular noise distribution regular-$\mathcal{HW}_{R_2, t_2}$ is $\log(\frac{N_2^{t_2}}{t_2!})$, to settle for the same level of security we only need $t_1 = t_2$ and $N_1 = N_2/e$ where e is the natural logarithm. Placing this back to table 3, the total FullEval time (i.e., the seed expanding time of PCG) can have an extra $\times$e speedup.

Yaxin: Double check i the cost of practical attacks against R^c-LPN scales with entropy.

Next we plug in concrete parameters and evaluate the performance of different DMPF schemes under different PCG parameter settings.

Define the number of noisy coordinate $W := ct$. We set (λ, c, N, W) such that the best attack requires at least 2^λ arithmetic operations over field $\mathbb{F}_p$ of size approximately 2^{128}. According to [3], for R from an irreducible F, we lowerbound the number of arithmetic operations by $N \cdot (c \cdot \frac{N}{N-1})^W \approx N \cdot c^W$. Yaxin: TBD: Check [11] for latest update about this cost. For R from a reducible F, we lowerbound the number of arithmetic operations by $2^i \cdot c^{W_i}$, where $i := \underset{1 \leq i \leq \log N}{\arg\min} \left(2^i \cdot c^{W_i}\right)$ and $W_i := W - cn + (c(n-1) + W) \cdot \left(1 - \frac{1}{n}\right)^{W/c-1}$. In the end we round W such that $t = W/c$ is an integer.

Yaxin: Previous calculation as a reference: choosing the big-state DMPF for $t < 8$ and the OKVS-DMPF for $t \geq 8$ gives at least $\times 2$ acceleration on expand time over other choices with sacrifice on the keysize. There is a tradeoff between the batch-code and OKVS-DMPF in that the OKVS-DMPF always provides a $\sim \times 2$ acceleration on expand time, but a loss in seed size that when t is large it may blow up the seed size to $\sim \times 2$ that of the batch-code-DMPF.

4.2 Unbalanced PSI-WCA

A private set intersection (PSI) protocol allows two parties with input X, Y being two sets to learn about their intersection $X \cap Y$ without revealing additional information of X or Y. We denote by PSI-WCA (weighted cardinality) a variant of PSI that computes the weighted cardinality of elements in $X \cap Y$ where the weights are determined by a pre-fixed function $w(\cdot)$.

We will be interested in *unbalanced* PSI-WCA where $|X| \gg |Y|$ and the output should be received by the party holding Y. In this problem we call the party holding X as the server, and the party holding Y as the client. If further the big set X is held by two non-colluding servers, then such an unbalanced PSI-WCA protocol can be constructed from DMPF, as suggested in [7]:

- The client invokes DMPF.Gen$(1^\lambda, \hat{f}_{Y, w(Y)}) \rightarrow (k_0, k_1)$, where $w(Y)$ is the set of weights of elements in Y. Then the client send k_0 to server 0 and k_1 to server 1.
- Server b computes $s_b = \sum_{x \in X}$ DMPF.Eval$_b(1^\lambda, k_b, x)$ and send it back to the client.
- The client computes $s_0 + s_1$, which will be the weighted cardinality of $X \cap Y$.

One caveat is that this protocol reveals information about Y that is leaked by DMPF. Plugging in any DMPF instantiations we have mentioned, the size of $|Y|$ will be leaked to the servers.

The cost of our unbalanced PSI-WCA can be computed as follows:

- The communication cost equals the keysize of DMPF.
- The client computation time equals the key generation time of DMPF.
- The server computation time equals $|X| \times$ the evaluation time of DMPF.

We'll instantiate DMPF in different ways as listed in table 1. As suggested in [7], we take an infeasibly large domain for the sets X and Y to locate, whose size is $N = 2^{128}$. The set sizes $|X|$ and $|Y|$ can vary depending on application scenarios. Since $|Y|$ is the crucial factor that distinguishes different DMPF instantiations, we will only consider the change of $|Y|$. The costs of PSI-WCA protocols under different settings of $|Y|$ are listed in table 4.

Yaxin: Previous calculation as a reference: A short conclusion is using big-state DMPF for $t < 64$ and the OKVS-DMPF for $t \geq 64$ gives at $\sim \times 2$ faster Eval() time and faster Gen() time compared to the naive and batch-code construction. The keysize ($\propto$ communication complexity) of our choice is usually smaller than the batch-code DMPF and slightly larger than the naive construction.

Table 3: Seed size and expanding time of PCG protocols for the same (λ, N, c, t) with different choices of noise distributions in module-LPN assumption, and with different DMPF instantiations. We use construction 4 as an instantiation of OKVS. The seed size is represented by total DMPF share size and the expanding time is represented by total DMPF.FullEval time. The PRG evaluations in the first $\log N$ layers and in the convert layer are both regarded as the same PRG. e in the second row represents the expansion parameter for PBC, and e' in the last row represents the expansion parameter for OKVS.

DMPF instantiation	Noise type	Total share size	Total FullEval time (only listed PRG and OKVS)
Sum of DPFs	regular	$c^2t^2\lambda\log(2N/t) + c^2t^2\log p$	$4c^2tN\times$PRG
	nonregular	$c^2t^2\lambda\log(2N) + c^2t^2\log p$	$4c^2t^2N\times$PRG
Batch-code DMPF	regular	$ec^2t^2\lambda\log(\frac{wN}{et}) + ec^2t^2\log p$	$8c^2wN\times$PRG
	nonregular	$ec^2t^2\lambda\log(\frac{2wN}{et^2}) + ec^2t^2\log p$	$4c^2wN\times$PRG
Big-state DMPF	regular	$c^2t^2(\lambda + \frac{4}{3}t)\log(2N) + c^2t^2\log p$	$8c^2N\times$PRG[*1]
	nonregular	$c^2t^2(\lambda + 2t)\log(2N) + c^2t^2\log p$	$4c^2N\times$PRG[*]
OKVS-based DMPF	regular	$e'c^2t^2\lambda\log(2N/t) + e'c^2t^2\log p$	$8c^2N\times$PRG$+8c^2N\times$OKVS.Decode
	nonregular	$e'c^2t^2\lambda\log(2N) + e'c^2t^2\log p$	$4c^2N\times$PRG$+4c^2N\times$OKVS.Decode

[1] The PRG used in big-state DMPF maps from $\{0,1\}^\lambda$ to $\{0,1\}^{2\lambda+2t^2}$ whose computation time should grow with t^2.

Table 4: Communication cost, client and server computation time of the PSI-WCA protocol for domain size $N = 2^{128}$, weight group $\mathbb{G}$, and different choices of client's set size $|Y|$. We use construction 4 as an instantiation of OKVS. The PRG evaluations in the first $\log N$ layers and in the convert layer are both regarded as the same PRG. e in the second row represents the expansion parameter for PBC, and e' in the last row represents the expansion parameter for OKVS.

DMPF instantiation	Communication cost	Client computation time	Server computation time
Sum of DPFs	$|Y|\lambda\log N + |Y|\log|\mathbb{G}|$	$2|Y|\log N\times$PRG	$|X|\cdot|Y|\log N\times$PRG
Batch-code DMPF	$e|Y|\lambda\log(\frac{wN}{e|Y|}) + e|Y|\log|\mathbb{G}|$	$2e|Y|\log(\frac{wN}{e|Y|})\times$PRG	$w|X|\log(\frac{wN}{e|Y|})\times$PRG
Big-state DMPF	$|Y|(\lambda + 2|Y|)\log N + |Y|\log|\mathbb{G}|$	$2|Y|\log N\times$PRG[*1]	$|X|\log N\times$PRG[*]
OKVS-based DMPF	$e'|Y|\lambda\log N + e'|Y|\log|\mathbb{G}|$	$2|Y|\log N\times$PRG$+\log N\times$OKVS.Encode	$|X|(\log N\times$PRG$+\log N\times$OKVS.Decode$)$

[1] The PRG used in big-state DMPF maps from $\{0,1\}^\lambda$ to $\{0,1\}^{2\lambda+2|Y|}$ whose computation time should grow with $|Y|$.

4.3 Security analysis

4.4 Heavy-hitters

private heavy-hitter
or parallel ORAM?

5 ACKNOWLEDGMENTS

tbd

REFERENCES

[1] Sebastian Angel, Hao Chen, Kim Laine, and Srinath Setty. 2017. PIR with compressed queries and amortized query processing. Cryptology ePrint Archive, Paper 2017/1142. https://eprint.iacr.org/2017/1142 https://eprint.iacr.org/2017/1142.
[2] Elette Boyle, Geoffroy Couteau, Niv Gilboa, and Yuval Ishai. 2019. Compressing Vector OLE. Cryptology ePrint Archive, Paper 2019/273. https://doi.org/10.1145/3243734.3243868 https://eprint.iacr.org/2019/273.
[3] Elette Boyle, Geoffroy Couteau, Niv Gilboa, Yuval Ishai, Lisa Kohl, and Peter Scholl. 2022. Efficient Pseudorandom Correlation Generators from Ring-LPN.

Cryptology ePrint Archive, Paper 2022/1035. https://doi.org/10.1007/978-3-030-56880-1_14 https://eprint.iacr.org/2022/1035.
[4] Elette Boyle, Niv Gilboa, and Yuval Ishai. 2018. Function Secret Sharing: Improvements and Extensions. Cryptology ePrint Archive, Paper 2018/707. https://eprint.iacr.org/2018/707 https://eprint.iacr.org/2018/707.
[5] Hao Chen, Kim Laine, and Peter Rindal. 2017. Fast Private Set Intersection from Homomorphic Encryption. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS '17)*. Association for Computing Machinery, New York, NY, USA, 1243–1255. https://doi.org/10.1145/3133956.3134061
[6] Leo de Castro and Antigoni Polychroniadou. 2021. Lightweight, Maliciously Secure Verifiable Function Secret Sharing. Cryptology ePrint Archive, Paper 2021/580. https://eprint.iacr.org/2021/580 https://eprint.iacr.org/2021/580.
[7] Samuel Dittmer, Yuval Ishai, Steve Lu, Rafail Ostrovsky, Mohamed Elsabagh, Nikolaos Kiourtis, Brian Schulte, and Angelos Stavrou. 2020. Function Secret Sharing for PSI-CA: With Applications to Private Contact Tracing. Cryptology ePrint Archive, Paper 2020/1599. https://eprint.iacr.org/2020/1599 https://eprint.iacr.org/2020/1599.
[8] Gayathri Garimella, Benny Pinkas, Mike Rosulek, Ni Trieu, and Avishay Yanai. 2021. Oblivious Key-Value Stores and Amplification for Private Set Intersection. Cryptology ePrint Archive, Paper 2021/883. https://eprint.iacr.org/2021/883 https://eprint.iacr.org/2021/883.

[9] Niv Gilboa and Yuval Ishai. 2014. Distributed Point Functions and Their Applications. In *Advances in Cryptology – EUROCRYPT 2014*, Phong Q. Nguyen and Elisabeth Oswald (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 640–658.

[10] Yuval Ishai, Eyal Kushilevitz, Rafail Ostrovsky, and Amit Sahai. 2004. Batch Codes and Their Applications. In *Proceedings of the Thirty-Sixth Annual ACM Symposium on Theory of Computing* (Chicago, IL, USA) *(STOC '04)*. Association for Computing Machinery, New York, NY, USA, 262–271. https://doi.org/10.1145/1007352.1007396

[11] Hanlin Liu, Xiao Wang, Kang Yang, and Yu Yu. 2022. The Hardness of LPN over Any Integer Ring and Field for PCG Applications. Cryptology ePrint Archive, Paper 2022/712. https://eprint.iacr.org/2022/712 https://eprint.iacr.org/2022/712.

[12] Rasmus Pagh and Flemming Friche Rodler. 2001. Cuckoo Hashing. In *Algorithms — ESA 2001*, Friedhelm Meyer auf der Heide (Ed.). Springer Berlin Heidelberg,

Berlin, Heidelberg, 121–133.

[13] Srinivasan Raghuraman and Peter Rindal. 2022. Blazing Fast PSI from Improved OKVS and Subfield VOLE. Cryptology ePrint Archive, Paper 2022/320. https://eprint.iacr.org/2022/320 https://eprint.iacr.org/2022/320.

[14] Phillipp Schoppmann, Adrià Gascón, Leonie Reichert, and Mariana Raykova. 2019. Distributed Vector-OLE: Improved Constructions and Implementation. Cryptology ePrint Archive, Paper 2019/1084. https://doi.org/10.1145/3319535.3363228 https://eprint.iacr.org/2019/1084.

A SECURITY PROOFS