

The Name of the Title Is Hope

tbd

ABSTRACT

tbd.

CCS CONCEPTS

• Theory of computation → Cryptographic primitives.

KEYWORDS

tbd

ACM Reference Format:

tbd. tbd. The Name of the Title Is Hope. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 3 pages. <https://doi.org/tbd>

1 INTRODUCTION

tbd

2 PRELIMINARY

2.1 Basic Notations

Point and multi-point functions. Given a domain size N and Abelian group \mathbb{G} , a point function $f_{\alpha,\beta} : [N] \rightarrow \mathbb{G}$ for $\alpha \in [N]$ and $\beta \in \mathbb{G}$ evaluates to β on input α and to $0 \in \mathbb{G}$ on all other inputs. We denote by $\hat{f}_{\alpha,\beta} = (N, \mathbb{G}, \alpha, \beta)$ the representation of such a point function. A multi-point function $f_{A,B} : [N] \rightarrow \mathbb{G}$ for $A = (\alpha_1, \dots, \alpha_t) \in [N]^t$ and $B = (\beta_1, \dots, \beta_t) \in \mathbb{G}^t$ evaluates to β_i on input α_i for $1 \leq i \leq t$ and to 0 on all other inputs. Denote $\hat{f}_{A,B}(N, \mathbb{G}, A, B)$ the representation of such a point function.

Enote: MPF. Also representation of groups.

2.2 Distributed Multi-Point Functions

Enote: should directly adapt to multi-point function case

We begin by defining a slightly generalized notion of distributed point functions (DPFs), which accounts for the extra parameter \mathbb{G}' .

DEFINITION 1 (DPF [1, 3]). A (2-party) distributed point function (DPF) is a triple of algorithms $\Pi = (\text{Gen}, \text{Eval}_0, \text{Eval}_1)$ with the following syntax:

- $\text{Gen}(1^\lambda, \hat{f}_{\alpha,\beta}) \rightarrow (k_0, k_1)$: On input security parameter $\lambda \in \mathbb{N}$ and point function description $\hat{f}_{\alpha,\beta} = (N, \mathbb{G}, \alpha, \beta)$, the (randomized) key generation algorithm Gen returns a pair of keys $k_0, k_1 \in \{0, 1\}^*$. We assume that N and \mathbb{G} are determined by each key.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference acronym 'XX, tbd, tbd

© tbd Association for Computing Machinery.
ACM ISBN tbd...\$15.00
<https://doi.org/tbd>

- $\text{Eval}_i(k_i, x) \rightarrow y_i$: On input key $k_i \in \{0, 1\}^*$ and input $x \in [N]$ the (deterministic) evaluation algorithm of server i , Eval_i returns $y_i \in \mathbb{G}$.

We require Π to satisfy the following requirements:

- **Correctness:** For every λ , $\hat{f} = \hat{f}_{\alpha,\beta} = (N, \mathbb{G}, \alpha, \beta)$ such that $\beta \in \mathbb{G}$, and $x \in [N]$, if $(k_0, k_1) \leftarrow \text{Gen}(1^\lambda, \hat{f})$, then

$$\Pr \left[\sum_{i=0}^1 \text{Eval}_i(k_i, x) = f_{\alpha,\beta}(x) \right] = 1$$

- **Security:** Consider the following semantic security challenge experiment for corrupted server $i \in \{0, 1\}$:

- (1) The adversary produces two point function descriptions $(\hat{f}^0 = (N, \mathbb{G}, \alpha_0, \beta_0), \hat{f}^1 = (N, \mathbb{G}, \alpha_1, \beta_1)) \leftarrow \mathcal{A}(1^\lambda)$, where $\alpha_i \in [N]$ and $\beta_i \in \mathbb{G}$.
- (2) The challenger samples $b \leftarrow \{0, 1\}$ and $(k_0, k_1) \leftarrow \text{Gen}(1^\lambda, \hat{f}^b)$.
- (3) The adversary outputs a guess $b' \leftarrow \mathcal{A}(k_i)$.

Denote by $\text{Adv}(1^\lambda, \mathcal{A}, i) = \Pr[b = b'] - 1/2$ the advantage of \mathcal{A} in guessing b in the above experiment. For every non-uniform polynomial time adversary \mathcal{A} there exists a negligible function ν such that $\text{Adv}(1^\lambda, \mathcal{A}, i) \leq \nu(\lambda)$ for all $\lambda \in \mathbb{N}$.

We will also be interested in applying the evaluation algorithm on all inputs. Given a DPF $(\text{Gen}, \text{Eval}_0, \text{Eval}_1)$, we denote by FullEval_i an algorithm which computes Eval_i on every input x . Hence, FullEval_i receives only a key k_i as input.

2.3 Batch Codes

combinatorial/probabilistic batch codes, with cuckoo hashing a concrete instantiation

2.4 Oblivious Key-Value Stores

DEFINITION 2 (OKVS[2, 4]). An Oblivious Key-Value Stores (OKVS) scheme is a pair of randomized algorithms $(\text{Encode}_r, \text{Decode}_r)$ with respect to a statistical security parameter λ_{stat} and a computational security parameter λ , a randomness space $\{0, 1\}^k$, a key space \mathcal{K} , a value space \mathcal{V} , input length n and output length m . The algorithms are of the following syntax:

- $\text{Encode}_r(\{(k_1, v_1), (k_2, v_2), \dots, (k_n, v_n)\}) \rightarrow P$: On input n key-value pairs with distinct keys, the encode algorithm with randomness r in the randomness space outputs an encoding $P \in \mathcal{V}^m \cup \perp$.
- $\text{Decode}_r(P, k) \rightarrow v$: On input a (nonempty) encoding from \mathcal{V}^m and a key $k \in \mathcal{K}$, output a value v .

We require the scheme to satisfy

- **Correctness:** For every $S \in (\mathcal{K} \times \mathcal{V})^n$, $\Pr_{r \leftarrow \{0, 1\}^k} [\text{Decode}_r(S) = \perp] \leq 2^{-\lambda_{\text{stat}}}$.

- **Obliviousness:** For any distinct key sets $\{k_1^0, k_2^0, \dots, k_n^0\}$ and $\{k_1^1, k_2^1, \dots, k_n^1\}$ that are different, if they are paired with random values then their encodings are computationally indistinguishable, i.e.,

$$\{r, \text{Encode}_r(\{(k_1^0, v_1), \dots, (k_n^0, v_n)\})\}_{v_1, \dots, v_n \leftarrow \mathcal{V}, r \leftarrow \{0,1\}^\kappa} \\ \approx_c \{r, \text{Encode}_r(\{(k_1^1, v_1), \dots, (k_n^1, v_n)\})\}_{v_1, \dots, v_n \leftarrow \mathcal{V}, r \leftarrow \{0,1\}^\kappa}$$

One can obtain a linear OKVS if in addition require:

- **Linearity:** There exists a function family $\{\text{row}_r : \mathcal{K} \rightarrow \mathcal{V}^m\}_{r \in \{0,1\}^\kappa}$ such that $\text{Decode}_r(P, k) = \langle \text{row}_r(k), P \rangle$.

The Encode process for a linear OKVS is the process of sampling a random P from the set of solutions of the linear system $\{\langle \text{row}_r(k_i), P \rangle = v_i\}_{1 \leq i \leq n}$.

We evaluate an OKVS scheme by its encoding size (output length m), encoding time and decoding time. We stress the following two (linear) OKVS constructions:

CONSTRUCTION 1 (POLYNOMIAL). Suppose $\mathcal{K} = \mathcal{V} = \mathbb{F}$ is a field. Set

- $\text{Encode}(\{(k_i, v_i)\}_{1 \leq i \leq n}) \rightarrow P$ where P is the coefficients of a $(n-1)$ -degree \mathbb{F} -polynomial g_P that $g_P(k_i) = v_i$ for $1 \leq i \leq n$.
- $\text{Decode}(P, k) \rightarrow g_P(k)$.

The polynomial OKVS possesses an optimal encoding size $m = n$, but the Encode process is a polynomial interpolation which is only known to be achieved in time $O(n \log^2 n)$. The time for a single decoding is $O(n)$ and that for batched decodings is (amortized) $O(\log^2 n)$.

An alternative construction that has near optimal encoding size but much better running time is as follows.

CONSTRUCTION 2 (3-HASH GARBLED CUCKOO TABLE (3H-GCT)[2, 4]). Suppose $\mathcal{V} = \mathbb{F}$ is a field. Set $\text{row}_r(k) := \text{row}_r^{\text{sparse}}(k) || \text{row}_r^{\text{dense}}(k)$ where $\text{row}_r^{\text{sparse}}$ outputs a uniformly random weight- w vector in $\{0, 1\}^{m_1}$, and $\text{row}_r^{\text{dense}}(k)$ outputs a short dense vector in \mathbb{F}^{m_2} .

- $\text{Encode}(\{(k_i, v_i)\}_{1 \leq i \leq n}) \rightarrow P$ where P is solved from the system $\{\langle \text{row}_r(k_i), P \rangle = v_i\}_{1 \leq i \leq n}$ using the triangulation algorithm in [4].
- $\text{Decode}(P, k) \rightarrow \langle \text{row}_r(k), P \rangle$.

This OKVS construction features a linear encoding time, constant decoding time while having a linear encoding size.

TBD: Carefully(!) recompute the comparison table for OKVS and insert

We take $w = 3$, the most common option that outruns other choices of w in terms of running time. Restating the conclusion in [4]: given n and λ_{stat} , the choices of e and \hat{g} are $e = 1.223 + \frac{\lambda_{\text{stat}} + 9.2}{4.144n^{0.55}}$ and $\hat{g} = \frac{\lambda_{\text{stat}}}{\log_2(en)}$.

TBD: mention some connections to cuckoo hashing

3 NEW DMPF CONSTRUCTIONS

3.1 Big-State DMPF

display the big-state DMPF (plus distributed gen)

3.2 Batch-Code DMPF

display the batch-code DMPF

3.3 OKVS-based DMPF

display the OKVS-based DMPF (plus distributed gen)

3.4 Comparison

Comparison table dependent to PRG & F-MUL(list the formulas?)

analyze tradeoff

distributed gen advantage

4 APPLICATIONS

4.1 PCG for OLE from Ring-LPN

Characterize parameters

show nonregular optimization

plug in new DMPF and show overall optimization

4.2 PSI-WCA

plug in new DMPF and analyze advantage interval

plug in distributed gen

4.3 Heavy-hitters

private heavy-hitter

or parallel ORAM?

5 ACKNOWLEDGMENTS

tbd

REFERENCES

- [1] Elette Boyle, Niv Gilboa, and Yuval Ishai. 2018. Function Secret Sharing: Improvements and Extensions. Cryptology ePrint Archive, Paper 2018/707. <https://eprint.iacr.org/2018/707> <https://eprint.iacr.org/2018/707>.
- [2] Gayathri Garimella, Benny Pinkas, Mike Rosulek, Ni Trieu, and Avishay Yanai. 2021. Oblivious Key-Value Stores and Amplification for Private Set Intersection. Cryptology ePrint Archive, Paper 2021/883. <https://eprint.iacr.org/2021/883> <https://eprint.iacr.org/2021/883>.
- [3] Niv Gilboa and Yuval Ishai. 2014. Distributed Point Functions and Their Applications. In *Advances in Cryptology – EUROCRYPT 2014*, Phong Q. Nguyen and Elisabeth Oswald (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 640–658.
- [4] Srinivasan Raghuraman and Peter Rindal. 2022. Blazing Fast PSI from Improved OKVS and Subfield VOLE. Cryptology ePrint Archive, Paper 2022/320. <https://eprint.iacr.org/2022/320> <https://eprint.iacr.org/2022/320>.

A BATCH-CODE DMPF SCHEME

B SECURITY PROOFS

Public parameters:

The multi-point function family $\{f_{A,B}\}$, an upperbound t of the number of nonzero points ($|A| \leq t$), input domain $[N] = \{0, 1\}^n$ and the output group \mathbb{G} .

Suppose there is a public PRG $G : \{0, 1\}^\lambda \rightarrow \{0, 1\}^{2\lambda+2l}$. Parse $G = G_0 || G_1$ to the left half and right half.

Suppose there is a public PRG $G_{\text{convert}} : \{0, 1\}^\lambda \rightarrow \mathbb{G}$.

procedure $\text{GEN}(1^\lambda, \hat{f}_{A,B})$

Denote $A = (\alpha_1, \dots, \alpha_t)$ in lexicographical order, $B = (\beta_1, \dots, \beta_t)$.

For $0 \leq i \leq n-1$, let $A^{(i)}$ denote the sorted and deduplicated list of i -bit prefixes of strings in A . Specifically, $A^{(0)} = [\epsilon]$.

For $0 \leq i \leq n-1$ and $b = 0, 1$, initialize empty lists $\text{seed}_b^{(i)}$ and $\text{sign}_b^{(i)}$.

Initialize($\{\text{seed}_b^{(0)}, \text{sign}_b^{(0)}\}_{b=0,1}$).

for $i = 1$ to n **do**

$CW^{(i)} \leftarrow \text{GenCW}(A, B, \{\text{seed}_b^{(i-1)}, \text{sign}_b^{(i-1)}\}_{b=0,1})$.

for $k = 1$ to $|A^{(i-1)}|$ and $z = 0, 1$ **do****if** $A^{(i-1)}[k] || z \in A^{(i)}$ **then**

For $b = 0, 1$, compute $\text{temp}_b \leftarrow \text{Correct}(A^{(i-1)}[k] || z, \text{seed}_b^{(i-1)}[k], \text{sign}_b^{(i-1)}[k], CW^{(i)})$.

Append the first λ bit of temp_b to $\text{seed}_b^{(i)}$ and the rest to $\text{sign}_b^{(i)}$.

end if**end for****end for**

$CW^{(n+1)} \leftarrow \text{GenConvCW}(A, B, \{\text{seed}_b^{(n)}, \text{sign}_b^{(n)}\}_{b=0,1})$.

Set $k_b \leftarrow (\text{seed}_b^{(0)}, \text{sign}_b^{(0)}, CW^{(1)}, CW^{(2)}, \dots, CW^{(n+1)})$.

return (k_0, k_1) .

end procedure**procedure** $\text{EVAL}_b(1^\lambda, k_b, x)$

Parse $k_b = ([\text{seed}^{(0)}], [\text{sign}^{(0)}], CW^{(1)}, CW^{(2)}, \dots, CW^{(n+1)})$.

Denote $x = x_1 x_2 \dots x_n$.

for $i = 1$ to n **do**

$\text{seed}^{(i)} || \text{sign}^{(i)} \leftarrow \text{Correct}(x_1 \dots x_i, \text{seed}^{(i-1)}, \text{sign}^{(i-1)}, CW^{(i)})$ where $\text{seed}^{(i)}$ is λ -bit.

end for

return $\text{ConvCorrect}(x, \text{seed}^{(n)}, \text{sign}^{(n)}, CW^{(n+1)})$.

end procedure

Figure 1: The paradigm of our DMPF schemes. We leave the PRG expand length l , methods Initialize, GenCW, GenConvCW, Correct, ConvCorrect to be determined by specific constructions.